


TensorFlow: A Framework for Scalable Machine Learning

Yuefeng Zhou, Google Brain



**You probably
want to know...**

What is TensorFlow?

Why did we create TensorFlow?

How does Tensorflow Work?

Example: Linear Regression

Example: Convolutional Neural Network

Distributed TensorFlow






Fast, flexible, and scalable
open-source machine learning
library

One system for research and
production

Runs on CPU, GPU, TPU, and
Mobile

Apache 2.0 license

TensorFlow Handles Complexity



Modeling complexity



**Distributed
System**



**Heterogenous
System**




Under the Hood

A multidimensional array.



TensorFlow



A graph of operations.

The TensorFlow Graph

Computation is defined as a graph

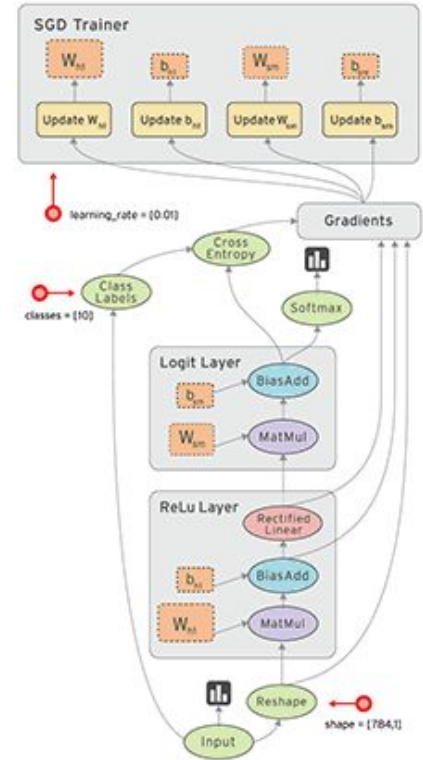
Graph is defined in high-level language (Python)

Graph is compiled and optimized

Graph is executed (in parts or fully) on available low level devices (CPU, GPU, TPU)

Nodes represent computations and state

Data (tensors) flow along edges



Build a graph; then run it.


...

```
c = tf.add(a, b)
```


...

```
session = tf.Session()
```

```
value_of_c = session.run(c, {a=1, b=2})
```

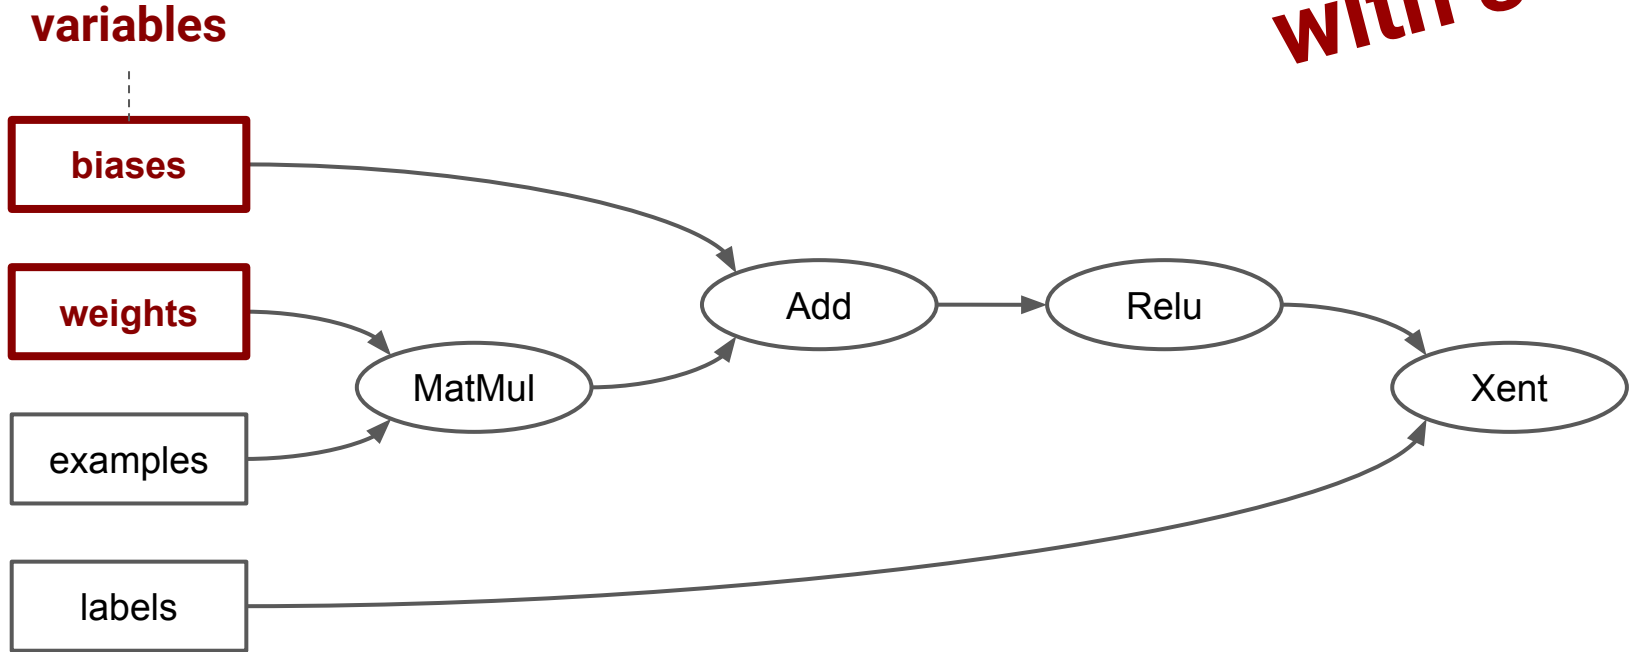


Any Computation is a TensorFlow Graph




Any Computation is a TensorFlow Graph

with state



Automatic Differentiation


**Automatically add ops which
compute gradients for variables**



Any Computation is a TensorFlow Graph


Simple gradient descent:

with state



Any Computation is a TensorFlow Graph

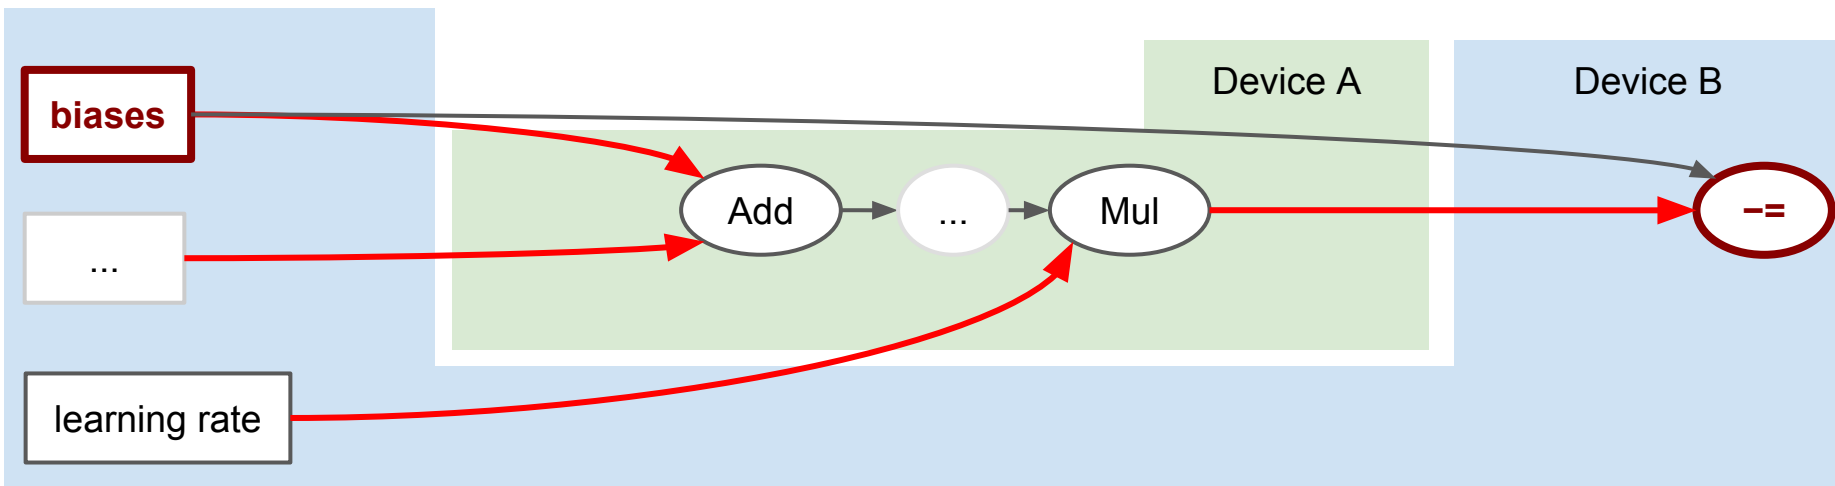
distributed



Devices: Processes, Machines, CPUs, GPUs, TPUs, etc

Send and Receive Nodes

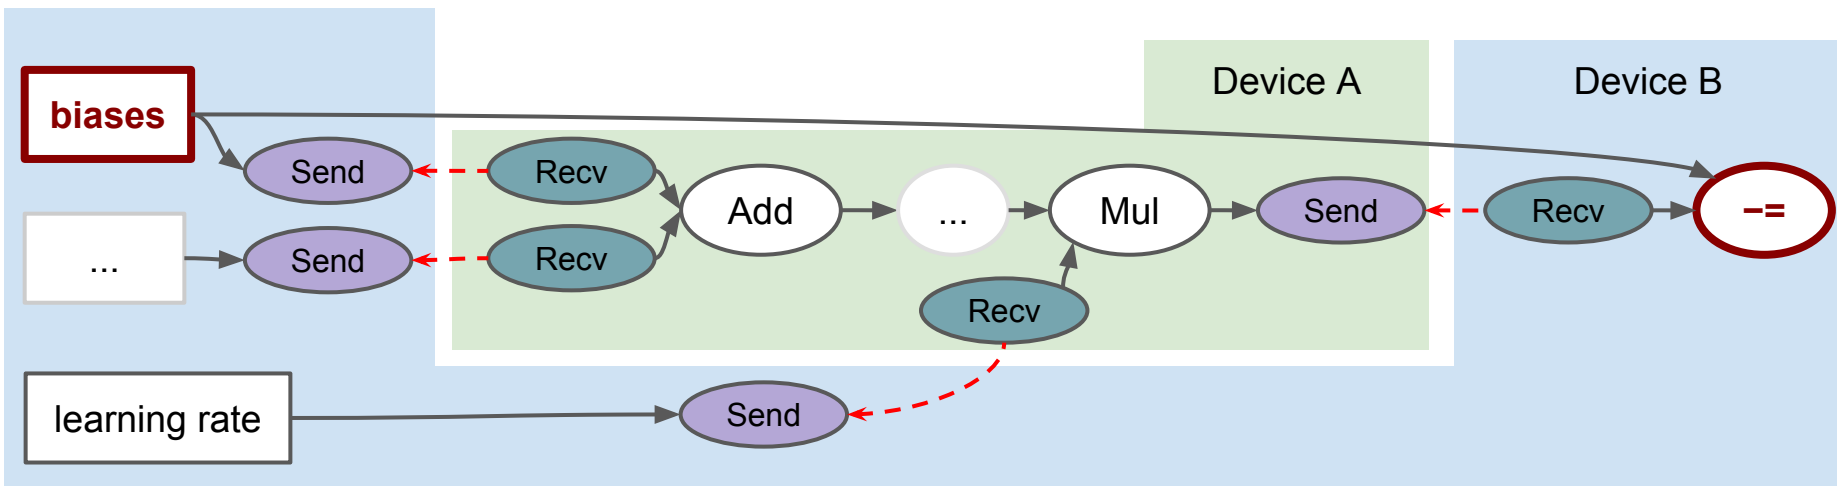
distributed



Devices: Processes, Machines, CPUs, GPUs, TPUs, etc

Send and Receive Nodes

distributed



Devices: Processes, Machines, CPUs, GPUs, TPUs, etc



Linear Regression

Linear Regression


result

input

$$y = Wx + b$$

parameters

The diagram shows the equation $y = Wx + b$. An arrow labeled 'result' points to y . An arrow labeled 'input' points to x . An arrow labeled 'parameters' points to both W and b .



What are we trying to do?

Mystery equation: $y = 0.1 * x + 0.3 + \text{noise}$

Model: $y = W * x + b$

Objective: Given enough (x, y) value samples, figure out the value of W and b .

$y = Wx + b$ in TensorFlow

```
import tensorflow as tf
```

$y = Wx + b$ in TensorFlow

```
import tensorflow as tf

x = tf.placeholder(shape=[None],
                   dtype=tf.float32, name="x")
```

$y = Wx + b$ in TensorFlow

```
import tensorflow as tf

x = tf.placeholder(shape=[None],
                    dtype=tf.float32, name="x")

W = tf.get_variable(shape=[], name="W")
```

$y = Wx + b$ in TensorFlow

```
import tensorflow as tf

x = tf.placeholder(shape=[None],
                   dtype=tf.float32, name="x")

W = tf.get_variable(shape=[], name="W")

b = tf.get_variable(shape=[], name="b")
```


$y = Wx + b$ in TensorFlow

```
import tensorflow as tf

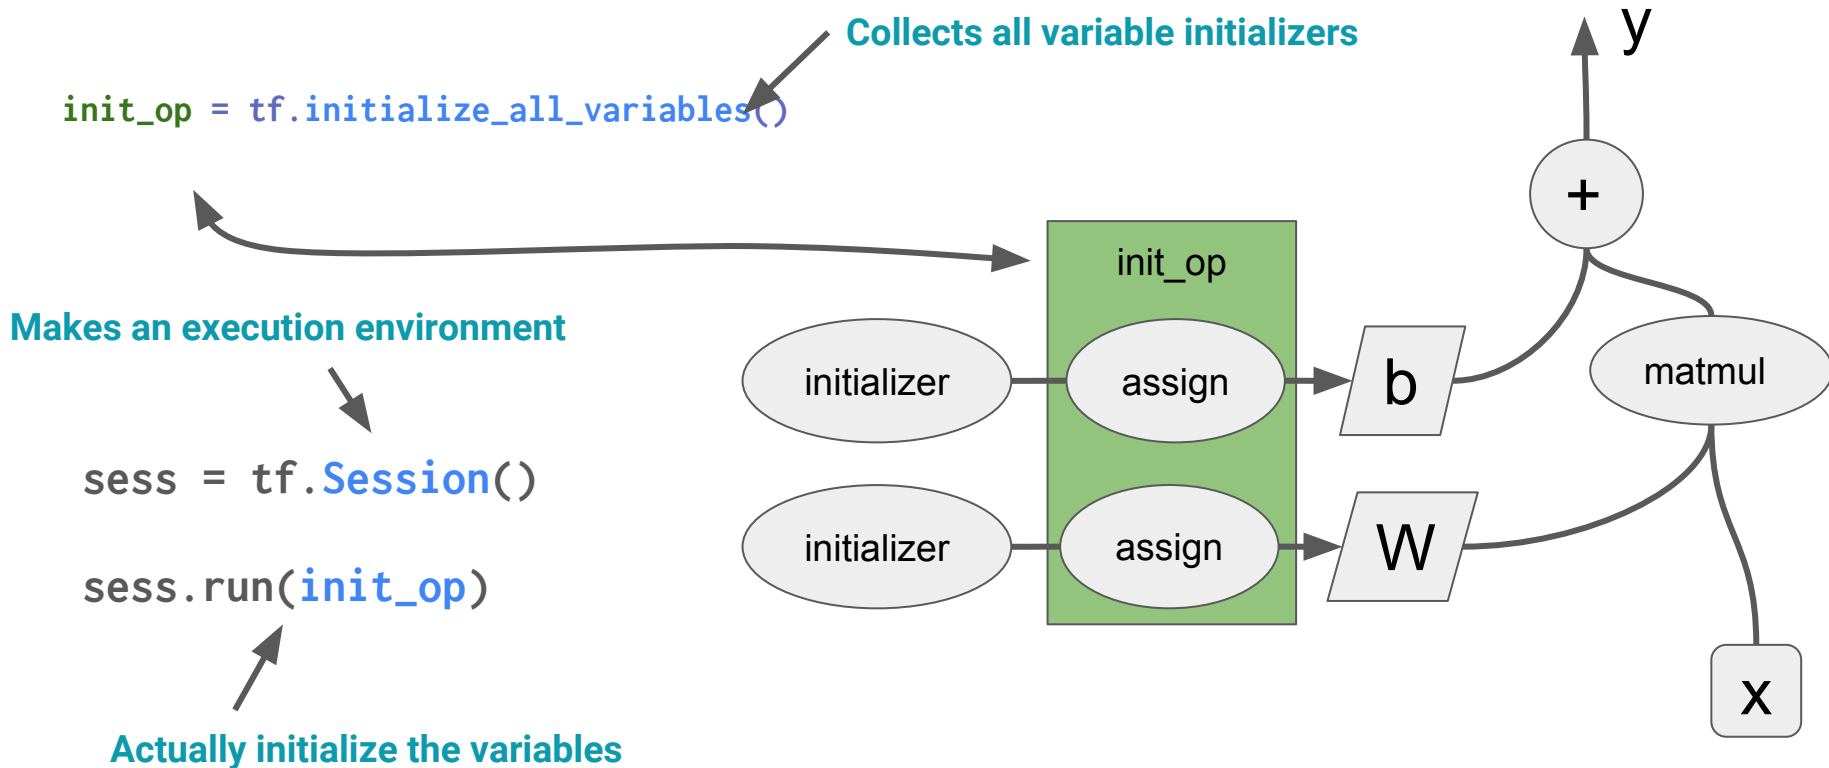
x = tf.placeholder(shape=[None],
                    dtype=tf.float32, name="x")

W = tf.get_variable(shape=[], name="W")
b = tf.get_variable(shape=[], name="b")

y = W * x + b
```



Variables Must be Initialized




Running the Computation

```
x_in = [3]
```

```
sess.run(y, feed_dict={x: x_in})
```

Only what's used to compute a fetch will be evaluated
All Tensors can be fed, but all placeholders must be fed



Putting it all together

```
import tensorflow as tf
x = tf.placeholder(shape=[None],
                   dtype=tf.float32,
                   name='x')

W = tf.get_variable(shape=[], name='W')
b = tf.get_variable(shape=[], name='b')
y = W * x + b
```

```
with tf.Session() as sess:
    sess.run(tf.initialize_all_variables())
    print(sess.run(y, feed_dict={x: x_in}))
```

} Build the graph

} Prepare execution environment

} Initialize variables

} Run the computation (usually often)


Define a Loss

Given y , y_{train} compute a loss, for instance:

$$L = (y - y_{train})^2$$

```
# create an operation that calculates loss.  
loss = tf.reduce_mean(tf.square(y - y_train))
```

Minimize loss: optimizers



`tf.train.AdadeltaOptimizer`
`tf.train.AdagradOptimizer`
`tf.train.AdagradDAOptimizer`
`tf.train.AdamOptimizer`
...

Train

Feed (x, y_{label}) pairs and adjust W and b to decrease the loss.

$$W \leftarrow W - \eta (dL/dW)$$


$$b \leftarrow b - \eta (dL/db)$$

```
# Create an optimizer
```

```
optimizer = tf.train.GradientDescentOptimizer(0.5)
```

```
# Create an operation that minimizes loss.
```

```
train = optimizer.minimize(loss)
```



TensorFlow computes
gradients automatically



Learning rate

Putting it all together

```
loss = tf.reduce_mean(tf.square(y - y_train))
```

```
optimizer = tf.train.GradientDescentOptimizer(0.5)
```

```
train = optimizer.minimize(loss)
```

```
with tf.Session() as sess:
```

```
    sess.run(tf.initialize_all_variables())
```

```
    for i in range(1000):
```

```
        sess.run(train, feed_dict={x: x_in, y_label: y_in})
```

} Define a loss

} Create an optimizer

} Op to minimize the loss

} Initialize variables

} Iteratively run the training op

Putting it all together

```
import tensorflow as tf

x = tf.placeholder(shape=[None], dtype=tf.float32, name='x')
W = tf.get_variable(shape=[], name='W')
b = tf.get_variable(shape=[], name='b')
y = W * x + b

loss = tf.reduce_mean(tf.square(y - y_train))

optimizer = tf.train.GradientDescentOptimizer(0.5)

train = optimizer.minimize(loss)

with tf.Session() as sess:

    sess.run(tf.initialize_all_variables())

    for i in range(1000):

        sess.run(train, feed_dict={x: x_in, y_label: y_in})
```

Build the graph

Define a loss

Create an optimizer

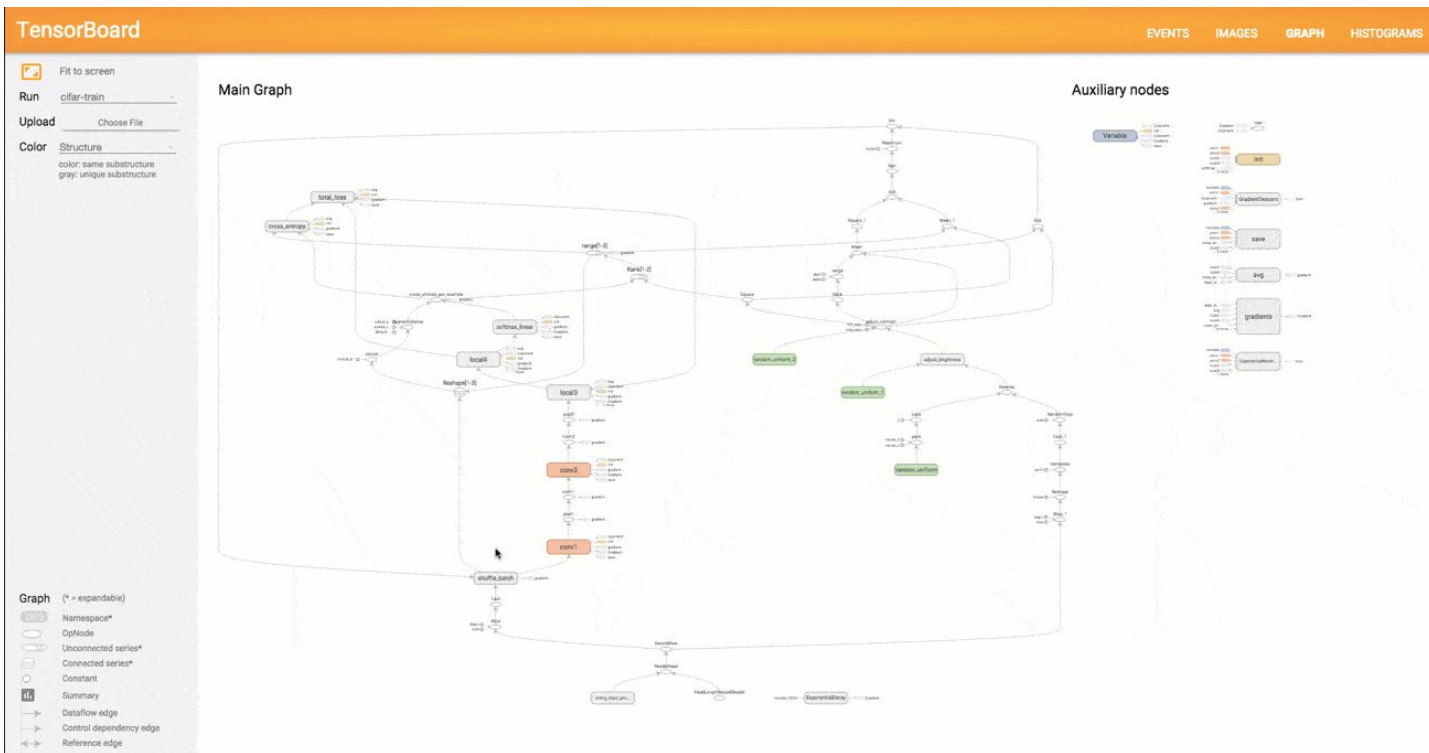
Op to minimize the
loss

Prepare environment

Initialize variables

Iteratively run the
training op

TensorBoard






Deep Convolutional Neural Network

IS THIS A
CAT or DOG?



CAT DOG



Remember linear regression?

```
import tensorflow as tf
x = tf.placeholder(shape=[None],
                   dtype=tf.float32,
                   name='x')

W = tf.get_variable(shape=[], name='W')
b = tf.get_variable(shape=[], name='b')
y = W * x + b

loss = tf.reduce_mean(tf.square(y - y_label))

optimizer = tf.train.GradientDescentOptimizer(0.5)


train = optimizer.minimize(loss)

...
```


Build the graph

Convolutional DNN

```
x = tf.contrib.layers.conv2d(x, kernel_size=[5,5], ...)
x = tf.contrib.layers.max_pool2d(x, kernel_size=[2,2], ...)
x = tf.contrib.layers.conv2d(x, kernel_size=[5,5], ...)
x = tf.contrib.layers.max_pool2d(x, kernel_size=[2,2], ...)
x = tf.contrib.layers.fully_connected(x, activation_fn=tf.nn.relu)
x = tf.contrib.layers.dropout(x, 0.5)
logits = tf.config.layers.linear(x)
```




Defining Complex Networks





Distributed TensorFlow

Data Parallelism (Between-Graph Replication)



Describe a cluster: ClusterSpec

```
tf.train.ClusterSpec({  
    "worker": [  
        "worker0.example.com:2222",  
        "worker1.example.com:2222",  
        "worker2.example.com:2222"  
    ],  
    "ps": [  
        "ps0.example.com:2222",  
        "ps1.example.com:2222"  
    ]  
})
```


Share the graph across devices

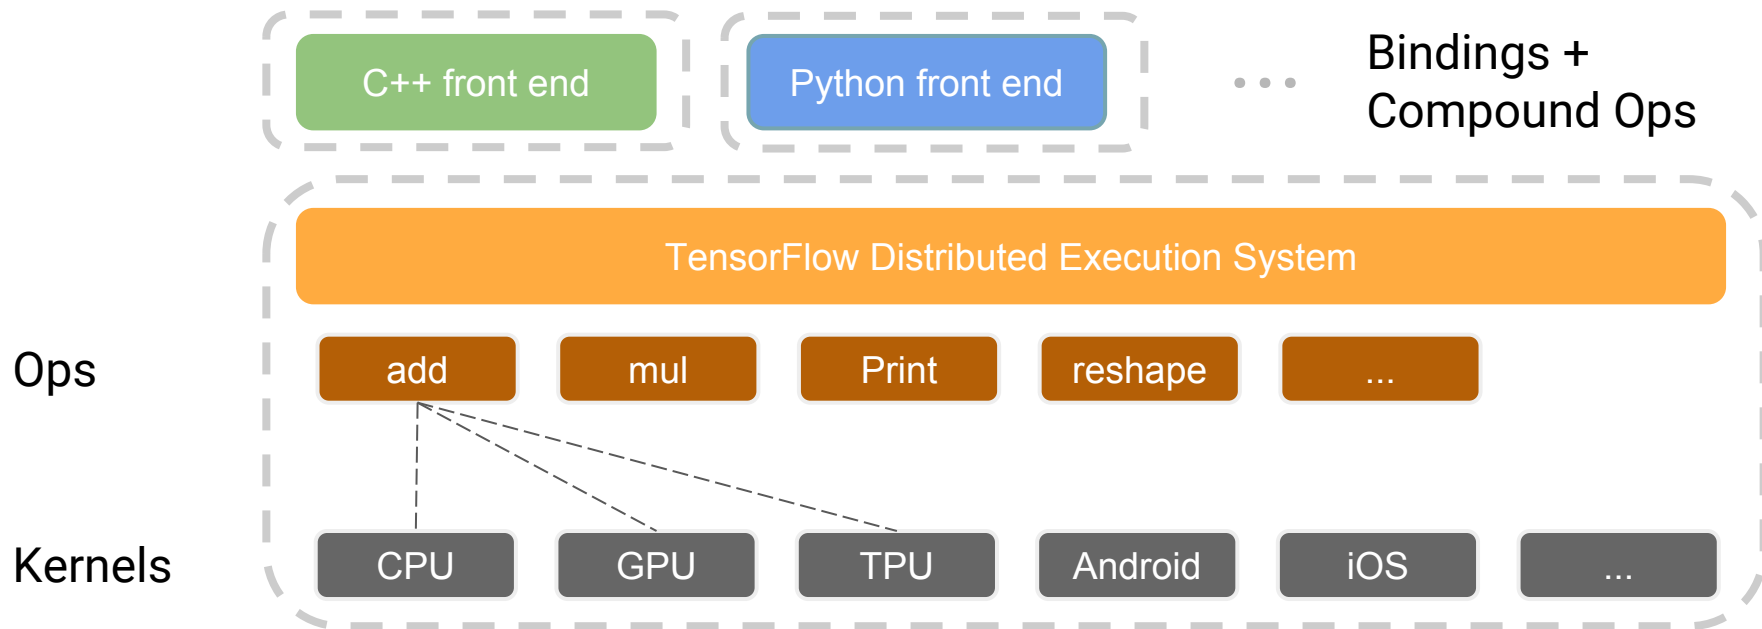
```
with tf.device("/job:ps/task:0"):
    weights_1 = tf.Variable(...)
    biases_1 = tf.Variable(...)
```

```
with tf.device("/job:ps/task:1"):
    weights_2 = tf.Variable(...)
    biases_2 = tf.Variable(...)
```

```
with tf.device("/job:worker/task:7"):
    input, labels = ...
    layer_1 = tf.nn.relu(tf.matmul(input, weights_1) + biases_1)
    logits = tf.nn.relu(tf.matmul(layer_1, weights_2) + biases_2)
    train_op = ...
```

```
with tf.Session("grpc://worker7.example.com:2222") as sess:
    for _ in range(10000):
        sess.run(train_op)
```

Architecture



Tutorials

Tutorials on [tensorflow.org](https://www.tensorflow.org)

Image recognition: https://www.tensorflow.org/tutorials/image_recognition

Word embeddings: <https://www.tensorflow.org/tutorials/word2vec>

Language Modeling: <https://www.tensorflow.org/tutorials/recurrent>

Translation: <https://www.tensorflow.org/tutorials/seq2seq>

Udacity Course: <https://www.udacity.com/course/deep-learning--ud730>



Q & A



Google Developer Day